



# Movie Recommendation System

Marissa F. , Vishal N., Alex Q., Neha T., Abbie A.

# Data Description & System Objective

- Dataset: export from IMDb
  - Distributed into tables on movie data and their ratings
  - ~ 20M rows of data
  - Films from 1894 - 2019
- Project objective: ingest large dataset using AWS and provide end users a recommendation
  - Intake selected preferences
  - Recommendation driven by ratings and machine learning.

## What to watch

### | Fan favorites


This week's top TV and movies



★ 7.0

The Platform

+ Watchlist



★ 8.6

Parasite

Watch options



★ 8.5

Money Heist

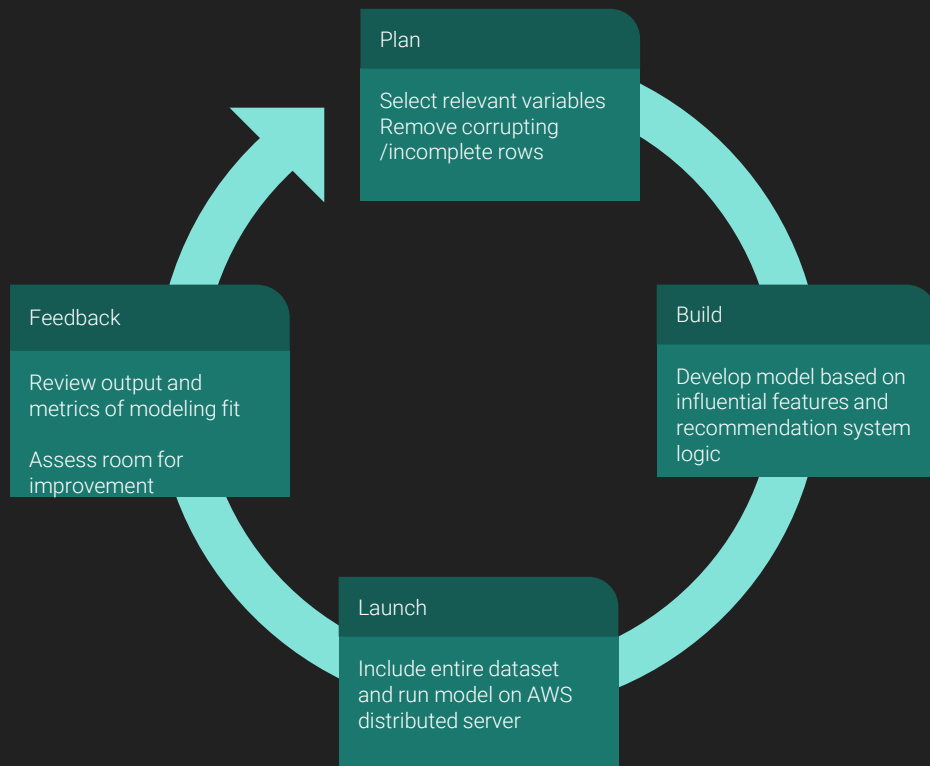
+ Watchlist

# Requirements

## Business Requirements and Attributes

Business Requirement ID	Functional/Technical	Business Requirement	Rationale	Source
1	Technical	Collect IMBD movie data.	This data will be used for predictive analysis.	Kaggle
2	Technical	Merge applicable tables to build single datasource.	This will allow for comprehensive analysis of all data.	Python
3	Functional	Populate single dataset to AWS	Store data in secure and faster environment	AWS
4	Technical	Wrangle data to remove duplicate entries, keeping only original language and title.	Remove duplicates that may weaken analysis	Python
5	Functional	Remove incomplete data , data with low number of votes for ratings (dependent variable), and more recent years only (as determined by quality of earlier movie data).	Ensure high quality data is used	Python
6	Functional	Perform basic statistical analysis to understand data.	Understand linkages and relationships between variables.	Python
7	Functional	Use data to make predictive recommendations on movies.	Utilize findings to recommend movies based on various preferences.	Python
8	Technical	Test model	Ensure model works	Python
9	Functional	Deploy model using AWS	Share with others!	AWS
10	Technical	Ensure application should shows horizontal scale	Fast computing speed	AWS

# Program Development



# Data Preparation

- Combine 4 tables, selecting relevant variables to include in models
- Convert csv into Pyspark dataframe
- Remove redundant entries
  - Redundancies caused by movies translated into various languages
- Delete entries without ratings
  - This is a key variable for modeling
- Export cleaned data into parquet files for consumption

Filtering data in PySpark for original titles, and removing rows with NA, filtering for movies with less than 250 votes (numVotes)

```
[ ] import pandas as pd
import numpy as np

[ ] from pyspark.sql.functions import lit, when, col, regexp_extract

[ ] movies_df_1 = movies_df.filter(movies_df["isOriginalTitle"] == '1')
movies_df_1.show()
```

titleId	ordering	title	region	language	types	attributes	isOriginalTitle	tconst	titleType
tt0000001	5	Carmencita	\N	\N	original	\N		1 tt0000001	short
tt0000002	1	Le clown et ses c...	\N	\N	original	\N		1 tt0000002	short Le clown
tt0000003	4	Pauvre Pierrot	\N	\N	original	\N		1 tt0000003	short Pa
tt0000004	1	Un bon bock	\N	\N	original	\N		1 tt0000004	short
tt0000005	8	Blacksmith Scene	\N	\N	original	\N		1 tt0000005	short Blac
tt0000006	5	Chinese Opium Den	\N	\N	original	\N		1 tt0000006	short Chine
tt0000007	2	Corbett and Court...	\N	\N	original	\N		1 tt0000007	short Corbett
tt0000008	4	Edison Kinetoscop...	\N	\N	original	\N		1 tt0000008	short Edison K
tt0000009	1	Miss Jerry	\N	\N	original	\N		1 tt0000009	movie
tt0000010	11	La sortie de l'us...	\N	\N	original	\N		1 tt0000010	short Exiting

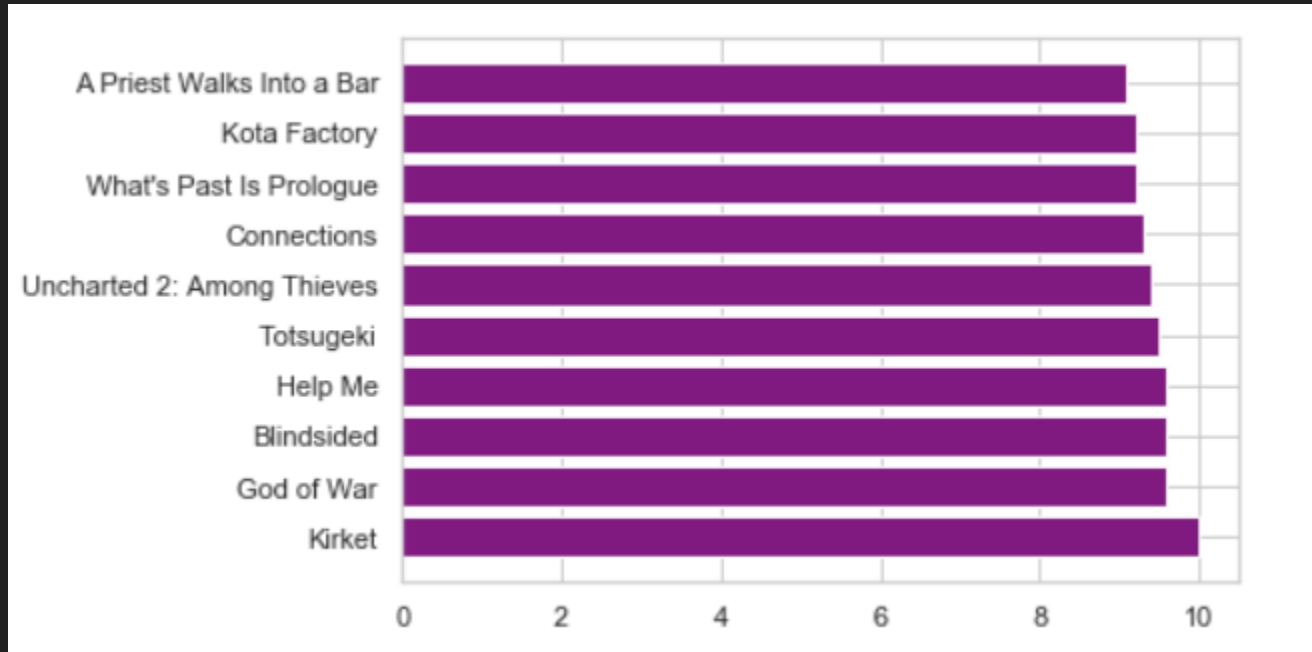
## Exporting code in PySpark

1. Parquet export code
2. CSV export code

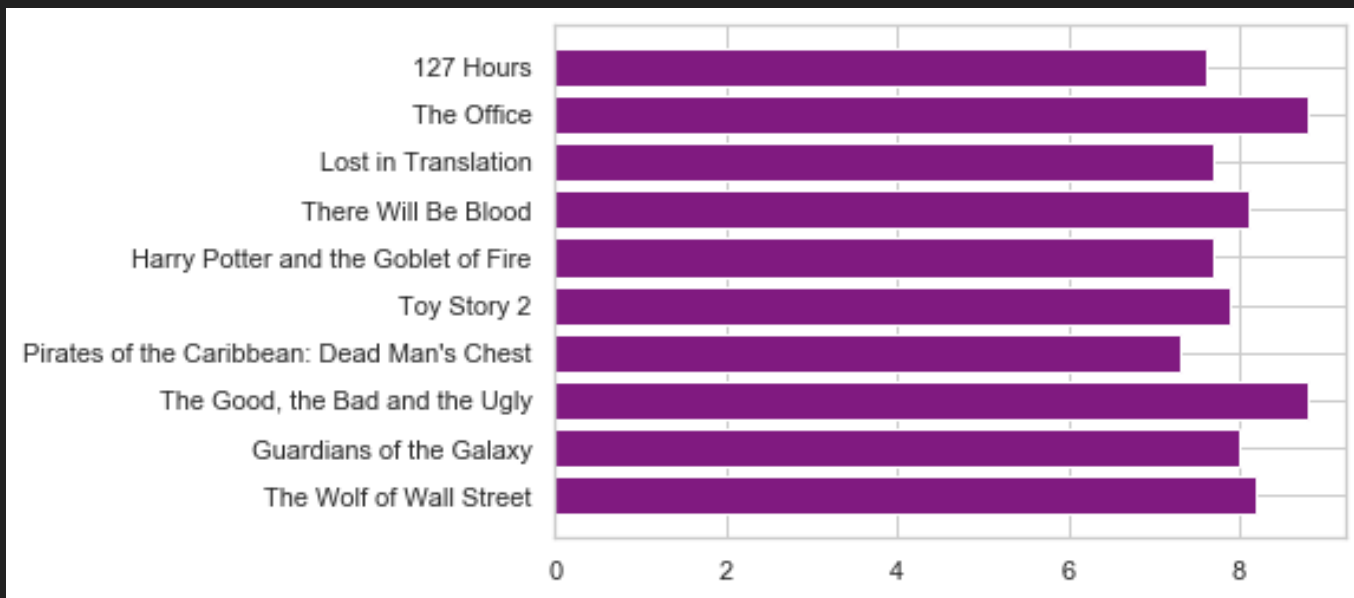
```
[ ] #partitioned file
movies_df_2.write.parquet("drive/My Drive/Unstructured Final Project/MoviesDatasetParquet")
#csv
movies_df_2.coalesce(1).write.mode("overwrite")\
.format("com.databricks.spark.csv")\
.option("header", "true")\
.option("sep", "|")\
.save('drive/My Drive/Unstructured Final Project/my_output_csv')
```

# Data Exploration & Analytics

# Highest Rated Content on IMDb

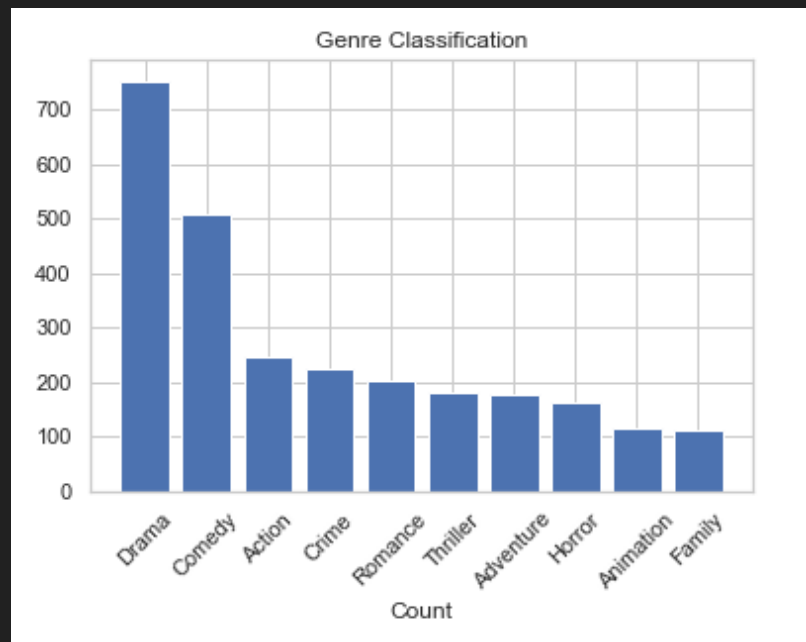


# Most Popular Content on IMDb

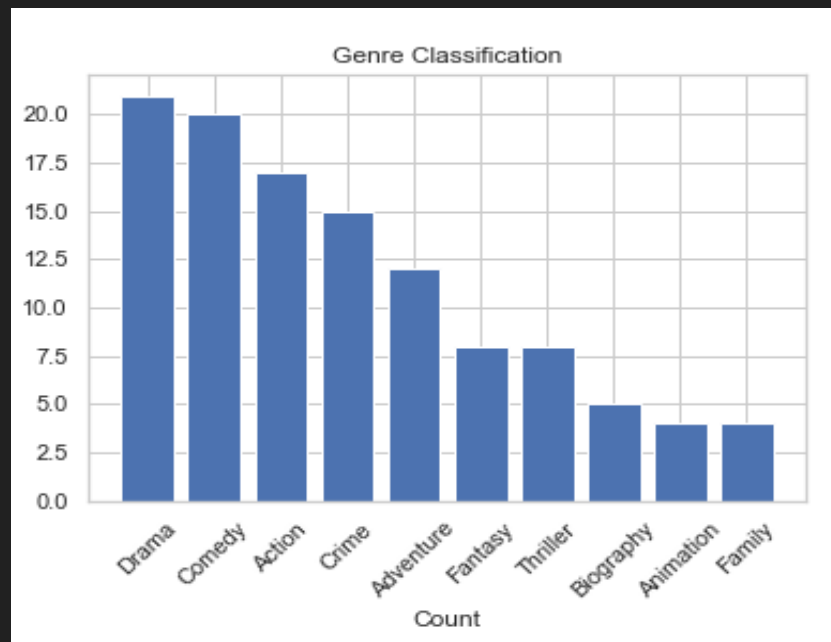




# Available Content



# Popular Content



# Machine Learning in PySpark to provide a recommendation for successful movies

## Predictor Variables:

1. Title Type: TV Series, movie, short film.
2. Number of Votes
3. Year
4. Movie Runtime
5. Genre

label	features	titleType	isAdult	startYear	runtimeMinutes	genres	numVotes	Success
0.0	(1330,[45,1188,13...]	short	0.0	1894	1	Documentary,Short	1550.0	0
0.0	(1330,[221,1188,1...]	short	0.0	1892	4	Animation,Comedy,...	1207.0	0
0.0	(1330,[22,1188,13...]	short	0.0	1893	1	Comedy,Short	1934.0	0
0.0	(1330,[616,1188,1...]	short	0.0	1894	1	Short,Sport	615.0	0
0.0	(1330,[45,1188,13...]	short	0.0	1894	1	Documentary,Short	1667.0	0
0.0	(1330,[45,1188,13...]	short	0.0	1895	1	Documentary,Short	5545.0	0
1.0	(1330,[785,1188,1...]	short	0.0	1896	1	Action,Documentar...	9435.0	1
0.0	(1330,[45,1188,13...]	short	0.0	1895	1	Documentary,Short	1447.0	0
1.0	(1330,[22,1188,13...]	short	0.0	1895	1	Comedy,Short	4111.0	1
0.0	(1330,[101,1188,1...]	short	0.0	1894	2	Animation,Short	741.0	0

only showing top 10 rows

## Approach:

### Output:


Predicts whether a movie can be classified as a success or not.

1. Set a threshold of an average rating  $\geq 6.4$  to determine whether the movie is successful.
2. Encode your feature variables (strings/categorical).
3. Transform your input variables into a vector of features that your model can work with.
4. Fit your model on the training sets and predict it on the test sets.

# Model Evaluation & Performance

titleType	numVotes	runTimeMinutes	startYear	genres	Success	prediction	probability
movie	457.0	136	2010	Drama	0	0.0	[0.67225970212625...
movie	323.0	88	2002	Drama	0	0.0	[0.74367873307500...
movie	430.0	79	2001	Drama	0	0.0	[0.72900975776502...
movie	2855.0	93	2001	Drama	0	0.0	[0.70419785400695...
movie	2577.0	97	2001	Drama	0	0.0	[0.69970649230874...
movie	337.0	108	2001	Drama	0	0.0	[0.69153701751024...
movie	888.0	116	2001	Drama	0	0.0	[0.67896847572049...
movie	493.0	89	2000	Drama	0	0.0	[0.71222581254970...
movie	2176.0	91	2000	Drama	0	0.0	[0.70479379287038...
movie	748.0	92	2000	Drama	0	0.0	[0.70758704334167...
movie	258.0	97	2000	Drama	0	0.0	[0.70241694037393...
movie	887.0	98	2000	Drama	0	0.0	[0.69926183413203...
movie	2048.0	98	2000	Drama	0	0.0	[0.69586695669028...
movie	269.0	99	2000	Drama	0	0.0	[0.69973173436651...
movie	2857.0	103	2000	Drama	0	0.0	[0.68671986477920...
movie	541.0	105	2000	Drama	0	0.0	[0.69088690434452...
movie	300.0	108	2000	Drama	0	0.0	[0.68753460196145...
movie	1014.0	115	2000	Drama	0	0.0	[0.67577226830311...
movie	307.0	139	2000	Drama	0	0.0	[0.64388864504795...
movie	302.0	55	1999	Drama	0	0.0	[0.77131689963880...
movie	1245.0	84	1999	Drama	0	0.0	[0.73479982404407...
movie	712.0	89	1999	Drama	0	0.0	[0.73004469720506...
movie	275.0	90	1999	Drama	0	0.0	[0.72999083882833...
movie	564.0	90	1999	Drama	0	0.0	[0.72920057599836...
movie	831.0	90	1999	Drama	0	0.0	[0.72846918075243...
movie	289.0	92	1999	Drama	0	0.0	[0.72744965878850...

True/Predicted	Positive	Negative	Total
Positive	6,341	3,207	16,451
Negative	2,594	10,110	5,801
Total	8,935	13,317	22,252



Metrics	
Sensitivity/Recall	66.4%
Specificity	79.6%
Precision	71.0%
Accuracy	73.9%

- Preferred model based on results: Logistic Regression from MLib in PySpark
- The outcome probabilities are shown for the combination of features in our data.
- The model has scope for improvement, either through the addition of more relevant predictor variables and additional data.

# Recommendation System

- Based on weighted ratings
  - $(v/(v+m) * R) + (m/(m+v) * C)$
  - $v$  = number of votes for the movie
  - $m$  = minimum number of votes to be considered (in Python, in the 10th percentile. In PySpark, over 100 votes)
  - $R$  = Average Rating of the movie
  - $C$  = Average Rating across all movies
- Three user inputs
  - Genre of choice
  - Type (i.e. movie, TV show, videogame)
  - Age

# Built in Python with a For Loop and Function

```
for i in range(0, len(movies)) :  
    v = movies.numVotes[i]  
    m = np.percentile(movies.averageRating, q = .10)  
    C = movies['averageRating'].mean()  
    R = movies.averageRating[i]  
    movies.WeightedRating[i] = (v/(v+m) * R) + (m/(m+v) * C)
```

```
def recommend(genrechoice, typechoice, agechoice = 0):  
    PossMovie = movies[movies['genres'].str.contains(genrechoice)]  
    PossMovie = PossMovie[PossMovie['titleType'].str.contains(typechoice)]  
    if (agechoice == "Old" or agechoice == 'old' or agechoice == "OLD"):  
        PossMovie = PossMovie[PossMovie.startYear <= 2000]  
    if (agechoice == "New" or agechoice == 'new' or agechoice == "NEW"):  
        PossMovie = PossMovie[PossMovie.startYear > 2000]  
    else:  
        PossMovie = PossMovie  
    Ordered = PossMovie.sort_values(by="WeightedRating", ascending = False)  
    Ordered = Ordered[['title', 'genres', 'titleType', 'startYear', 'WeightedRating']]  
    Top5 = Ordered[:5]  
    return Top5
```

```
recommend('Comedy', 'movie', 'New')
```

	title	genres	titleType	startYear	WeightedRating
45122	Shibu	Comedy	movie	2019	9.391544
59808	CM101MMXI Fundamentals	Comedy,Documentary	movie	2013	9.199900
67251	The Weight of Chains 2	Comedy,Documentary,History	movie	2014	8.996412
75873	10 Days Before the Wedding	Comedy,Drama,Musical	movie	2018	8.890196
34554	Anbe Sivam	Adventure,Comedy,Drama	movie	2003	8.799735

# Built in PySpark with a UDF and spark.sql

```
def weighted(v, m, C, R) :  
    return (v/(v+m) * R) + (m/(m+v) * C)
```

```
weighted_udf = udf(weighted, FloatType())
```

```
WeightedRating = weighted_udf(movies.numVotes, movies.MinimumVotes, movies.AverageTotalRating, movies.AverageRating)  
WeightedRating
```

```
Column<b'weighted(numVotes, MinimumVotes, AverageTotalRating, averageRating)'>
```

```
movies = movies.withColumn("WeightedRating", WeightedRating.cast(FloatType()))
```

```
movies = movies.select('title', 'titleType', 'startYear', 'genres', 'averageRating', 'WeightedRating')  
movies.show(10)
```

```
movies.createOrReplaceTempView("movies_sql")
```

```
recommendations_sql = spark.sql('''
    SELECT title, titleType, startYear, genres, WeightedRating
    FROM movies_sql
    WHERE genres LIKE '%Comedy%' AND titleType = 'movie' AND (startYear > 2000)
    ORDER BY WeightedRating desc
    LIMIT (5)
''')
```

```
recommendations_sql.show()
```

title	titleType	startYear	genres	WeightedRating
CM101MMXI Fundame...	movie	2013.0	Comedy,Documentary	9.193336
Shibu	movie	2019.0	Comedy	8.924001
Anbe Sivam	movie	2003.0	Adventure,Comedy,...	8.782449
The Weight of Cha...	movie	2014.0	Comedy,Documentar...	8.780598
Nuvvu Naaku Nachchav	movie	2001.0	Comedy,Family,Mus...	8.628696



# Findings & Conclusions

- Number of people voting on a movie is the most influential factor in a high scoring film
  - Popular films are perceived as good films
- The most popular movie genre is drama
- Deep categorical variables such as directors and actors could add value to the model via creating dummy columns for important figures in each area.
- Further Improvements for Deployment
  - Serialize the models and use a Flask API to generate a usable web app. A user could visit the website, select a genre, type, and the age of the film. The user would see the movie recommendations given by the model.